

MoE Expert Models: A Comprehensive Analysis of Three Specialized Fairness-Aware Neural Networks

Fairness Machine Learning Project

October 9, 2025

1 Problem Formulation

Let $\mathcal{C} = \{c_1, c_2, \dots, c_N\}$ denote a set of platforms (clients). Each platform c_i holds a private local dataset D_i that cannot be accessed by other platforms:

$$D_i = \{(x_{i,j}, s_{i,j}, y_{i,j})\}_{j=1}^{|D_i|}, \quad (1)$$

where $x_{i,j} \in \mathbb{R}^d$ is the feature vector of applicant j on platform i (including personal attributes and possibly representations of unstructured data), $s_{i,j} \in \{0, 1\}$ is a binary sensitive attribute (e.g., 0 = female, 1 = male), and $y_{i,j} \in \{0, 1\}$ is the target label indicating whether the applicant receives a positive decision.

We study fairness with respect to s and focus on both result fairness and procedural fairness. Let f_θ be a predictive model with parameters θ . For an input x , the model outputs class probabilities $p = f_\theta(x) \in [0, 1]^2$ and a predicted label $\hat{y} = \arg \max_c p_c$. Let $h_\theta(x)$ be the hidden representation and let $A_\theta(x) \in \mathbb{R}^d$ denote a normalized attribution ("attention") vector computed via a post-hoc interpreter (e.g., SHAP/IG/GradientSHAP), normalized so that $A_\theta(x) \geq 0$ and $\sum_{k=1}^d A_\theta(x)_k = 1$.

Result Fairness Metrics. We use Demographic Parity (DP) and Equal Opportunity (EO) gaps:

$$\text{DP} = |\mathbb{P}(\hat{y} = 1 \mid s = 1) - \mathbb{P}(\hat{y} = 1 \mid s = 0)|, \quad (2)$$

$$\text{EO} = |\mathbb{P}(\hat{y} = 1 \mid s = 1, y = 1) - \mathbb{P}(\hat{y} = 1 \mid s = 0, y = 1)|. \quad (3)$$

Procedural Fairness Metrics. We consider three complementary notions derived from explanations and attributions. Let $\text{EQ}(x) \in \mathbb{R}_{\geq 0}$ denote a scalar explanation quality for instance x produced by a fixed explainer. Define a top- K rule: let τ_K be the K -th percentile of $\{\text{EQ}(x_j)\}_{j=1}^n$ over the evaluation set and

$$\hat{q}_j = \mathbb{I}[\text{EQ}(x_j) \geq \tau_K]. \quad (4)$$

Let $G_s^K = \{j : s_j = s, \hat{q}_j = 1\}$ be the indices of top- K high-quality explanations within group s .

- **ΔREF** (Ratio-based Explanation Fairness): difference in high-quality rates between groups

$$\Delta\text{REF} = |\mathbb{P}(\hat{q} = 1 \mid s = 0) - \mathbb{P}(\hat{q} = 1 \mid s = 1)|. \quad (5)$$

Smaller is better; 0 indicates equal access to high-quality explanations.

- **ΔVEF** (Value-based Explanation Fairness): difference in average explanation quality among top- K within each group

$$\Delta\text{VEF} = \left| \frac{1}{|G_0^K|} \sum_{j \in G_0^K} \text{EQ}(x_j) - \frac{1}{|G_1^K|} \sum_{j \in G_1^K} \text{EQ}(x_j) \right|. \quad (6)$$

Smaller is better; 0 indicates equal quality among the top- K across groups.

- **ATT_JSD**: Jensen–Shannon divergence between the group-level attribution distributions

$$\text{ATT_JSD} = \text{JSD}(\bar{A}_0, \bar{A}_1) = \frac{1}{2} \text{KL}(\bar{A}_0 \| M) + \frac{1}{2} \text{KL}(\bar{A}_1 \| M), \quad (7)$$

$$M = \frac{1}{2}(\bar{A}_0 + \bar{A}_1). \quad (8)$$

Problem Definition. Given platforms \mathcal{C} with local datasets $\{D_i\}_{i=1}^N$, learn a global model f_θ that maximizes predictive utility while mitigating result fairness gaps (DP/EO) and procedural fairness disparities (REF/VEF/ATT_JSD), under data privacy constraints (no raw data sharing across platforms). A generic objective is

$$\min_{\theta} \mathbb{E}_{(x,y) \sim D} [\mathcal{L}_{\text{CE}}(f_\theta(x), y)] + \lambda_{dp} \text{DP} + \lambda_{eo} \text{EO} + \lambda_{ref} \Delta\text{REF} + \lambda_{vef} \Delta\text{VEF} + \lambda_{att} \text{ATT_JSD}, \quad (9)$$

where $D = \bigcup_{i=1}^N D_i$ conceptually denotes the union distribution (implemented in practice via local training/aggregation), and $\lambda_\bullet \geq 0$ control the utility–fairness trade-off. Lower values of DP/EO/REF/VEF/ATT_JSD indicate better fairness.

2 Introduction

This document provides a detailed mathematical and architectural analysis of the three expert models implemented in the Mixture of Experts (MoE) framework for fairness-aware machine learning. The system consists of three specialized neural network experts, each designed to address different aspects of fairness in machine learning: utility optimization, result-driven fairness, and procedural fairness.

3 Architecture Overview

All three experts inherit from a common **ExpertBase** class and share a fundamental Multi-Layer Perceptron (MLP) backbone. The base architecture consists of:

- Input layer: d_{input} dimensions
- Hidden layer: $d_{hidden} = 8$ dimensions (default)
- Output layer: 2 dimensions (binary classification)
- Activation: LeakyReLU
- Dropout: 0.3 (applied to hidden layer)

The forward pass of each expert follows the pattern:

$$h = \text{MLP}_{\text{hidden}}(x) \in \mathbb{R}^{d_{\text{hidden}}} \quad (10)$$

$$\ell = \text{MLP}_{\text{output}}(h) \in \mathbb{R}^2 \quad (11)$$

$$p = \text{softmax}(\ell) \in \mathbb{R}^2 \quad (12)$$

where h represents the hidden representation, ℓ the logits, and p the output probabilities.

4 Expert 1: Utility-Focused Expert

4.1 Objective

Expert 1 is designed to maximize predictive accuracy without explicit fairness constraints. It serves as a baseline utility-focused model.

4.2 Architecture

Expert 1 inherits directly from `ExpertBase` without additional components. The loss function is the standard cross-entropy loss:

$$\mathcal{L}_{E1} = \text{CrossEntropy}(p, y) = - \sum_{i=1}^n \sum_{c=0}^1 y_{i,c} \log(p_{i,c}) \quad (13)$$

where $y_{i,c}$ is the one-hot encoded true label for sample i and class c , and $p_{i,c}$ is the predicted probability.

4.3 Mathematical Formulation

The complete loss computation is:

$$\mathcal{L}_{E1} = \mathcal{L}_{CE} \quad (14)$$

$$\mathcal{L}_{CE} = -\frac{1}{n} \sum_{i=1}^n \log(p_{i,y_i}) \quad (15)$$

where y_i is the true class label for sample i .

5 Expert 2: Result-Driven Fairness Expert

5.1 Objective

Expert 2 addresses result-driven fairness by incorporating both representation alignment and demographic parity/equal opportunity constraints.

5.2 Architecture

Expert 2 extends `ExpertBase` with two additional hyperparameters:

- λ_{rep} : Weight for representation alignment loss
- λ_{fair} : Weight for fairness loss

5.3 Mathematical Formulation

The total loss function combines three components:

$$\mathcal{L}_{E2} = \mathcal{L}_{CE} + \lambda_{rep} \cdot \mathcal{L}_{rep} + \lambda_{fair} \cdot \mathcal{L}_{fair} \quad (16)$$

5.3.1 Representation Alignment Loss

The representation alignment loss minimizes the distance between group mean embeddings:

$$\mu_0 = \frac{1}{n_0} \sum_{i:s_i=0} h_i \quad (17)$$

$$\mu_1 = \frac{1}{n_1} \sum_{i:s_i=1} h_i \quad (18)$$

$$\mathcal{L}_{rep} = \|\mu_0 - \mu_1\|_2^2 \quad (19)$$

where $s_i \in \{0, 1\}$ is the sensitive attribute for sample i , and n_0, n_1 are the counts of samples in each group.

5.3.2 Fairness Loss

The fairness loss combines demographic parity (DP) and equal opportunity (EO) gaps:

$$\mathcal{L}_{fair} = \frac{\mathcal{L}_{DP} + \mathcal{L}_{EO}}{2} \quad (20)$$

Demographic Parity Gap:

$$p_0 = \frac{1}{n_0} \sum_{i:s_i=0} p_{i,1} \quad (21)$$

$$p_1 = \frac{1}{n_1} \sum_{i:s_i=1} p_{i,1} \quad (22)$$

$$\mathcal{L}_{DP} = |p_0 - p_1| \quad (23)$$

Equal Opportunity Gap:

$$p_0^{pos} = \frac{1}{n_0^{pos}} \sum_{i:s_i=0, y_i=1} p_{i,1} \quad (24)$$

$$p_1^{pos} = \frac{1}{n_1^{pos}} \sum_{i:s_i=1, y_i=1} p_{i,1} \quad (25)$$

$$\mathcal{L}_{EO} = |p_0^{pos} - p_1^{pos}| \quad (26)$$

where $p_{i,1}$ is the probability of positive class for sample i , and n_0^{pos}, n_1^{pos} are the counts of positive samples in each group.

6 Expert 3: Procedural Fairness Expert

6.1 Objective

Expert 3 addresses procedural fairness through attention alignment and adversarial debiasing techniques.

6.2 Architecture

Expert 3 extends **ExpertBase** with:

- $\lambda_{attention}$: Weight for attention fairness loss
- λ_{adv} : Weight for adversarial debiasing loss
- Gradient Reversal Layer (GRL) for adversarial training
- Two adversarial networks: one for hidden representations, one for logits

6.3 Mathematical Formulation

The total loss function is:

$$\mathcal{L}_{E3} = \mathcal{L}_{CE} + \lambda_{attention} \cdot \mathcal{L}_{attention} + \lambda_{adv} \cdot \mathcal{L}_{adv} \quad (27)$$

6.3.1 Attention Fairness Loss

The attention fairness loss uses Jensen-Shannon Divergence (JSD) to measure the difference in attention distributions between groups:

$$\mathcal{L}_{attention} = \text{JSD}(A_0, A_1) \quad (28)$$

$$\text{JSD}(A_0, A_1) = \frac{1}{2} \text{KL}(A_0 || M) + \frac{1}{2} \text{KL}(A_1 || M) \quad (29)$$

where $M = \frac{1}{2}(A_0 + A_1)$ is the average distribution, and A_0, A_1 are the mean attention weights for each group.

The attention weights are extracted using interpretability methods (SHAP, Integrated Gradients, LIME, or Gradient SHAP) and normalized:

$$A_{i,j} = \frac{|\text{attribution}_{i,j}|}{\sum_{k=1}^d |\text{attribution}_{i,k}| + \epsilon} \quad (30)$$

6.3.2 Adversarial Debiasing Loss

The adversarial loss prevents the model from learning sensitive attribute information:

$$\mathcal{L}_{adv} = \frac{\mathcal{L}_{adv}^{hidden} + \mathcal{L}_{adv}^{logits}}{2} \quad (31)$$

Hidden Representation Adversarial Loss:

$$h_{adv} = \text{GRL}(h) \quad (32)$$

$$\ell_{adv}^{hidden} = \text{MLP}_{adv}^{hidden}(h_{adv}) \quad (33)$$

$$\mathcal{L}_{adv}^{hidden} = \text{CrossEntropy}(\ell_{adv}^{hidden}, s) \quad (34)$$

Logits Adversarial Loss:

$$p_{adv} = \text{GRL}(p) \quad (35)$$

$$\ell_{adv}^{logits} = \text{MLP}_{adv}^{logits}(p_{adv}) \quad (36)$$

$$\mathcal{L}_{adv}^{logits} = \text{CrossEntropy}(\ell_{adv}^{logits}, s) \quad (37)$$

6.3.3 Gradient Reversal Layer

The GRL implements adversarial training by reversing gradients during backpropagation:

$$\text{GRL}(x) = x \quad (\text{forward pass}) \quad (38)$$

$$\frac{\partial \text{GRL}(x)}{\partial x} = -\lambda \quad (\text{backward pass}) \quad (39)$$

7 Implementation Details

7.1 MLP Backbone

The shared MLP backbone uses the following architecture:

- Linear layer: $d_{\text{input}} \rightarrow d_{\text{hidden}}$
- LeakyReLU activation
- Dropout (0.3)
- Linear layer: $d_{\text{hidden}} \rightarrow 2$

7.2 Adversarial Networks

Expert 3 includes two adversarial networks:

Hidden Adversarial Network:

- Input: d_{hidden} dimensions
- Hidden: $\max(4, d_{\text{hidden}}/2)$ dimensions
- Output: 2 dimensions (sensitive attribute prediction)

Logits Adversarial Network:

- Input: 2 dimensions (logits)
- Hidden: 8 dimensions
- Output: 2 dimensions (sensitive attribute prediction)

8 Training and Optimization

Each expert is trained independently with its specialized loss function. The training process involves:

1. Forward pass through the MLP backbone
2. Computation of specialized loss components
3. Backpropagation with appropriate gradient modifications (e.g., GRL for Expert 3)
4. Parameter updates using standard optimization algorithms

9 Gating Network Architecture

9.1 Overview

The Gating Network is responsible for selecting which expert to use for each input sample. It uses a REINFORCE-based policy gradient approach to learn optimal expert selection.

9.2 Architecture

The gating network consists of:

- Input layer: $d_{input} + 3 \cdot d_{classes} + 3 + 1$ dimensions
- Hidden layer: 16 dimensions (default)
- Output layer: 3 dimensions (one for each expert)
- Activation: ReLU
- Temperature parameter: $\tau = 1.0$ for softmax scaling

9.3 State Representation

The gating network receives an augmented state vector:

$$s = [x, p_1, p_2, p_3, \text{conf}_1, \text{conf}_2, \text{conf}_3, \text{disagree}] \quad (40)$$

where:

$$\text{conf}_i = \max(p_i) \quad (\text{confidence of expert } i) \quad (41)$$

$$\text{disagree} = \frac{1}{3} (\|p_1 - p_2\|_1 + \|p_1 - p_3\|_1 + \|p_2 - p_3\|_1) \quad (42)$$

9.4 Policy Formulation

The gating network outputs a probability distribution over experts:

$$\ell = \text{MLP}(s) \in \mathbb{R}^3 \quad (43)$$

$$\pi(a|s) = \text{softmax}(\ell/\tau) \in \mathbb{R}^3 \quad (44)$$

9.5 Action Sampling

Expert selection follows a categorical distribution:

$$a \sim \text{Categorical}(\pi(a|s)) \quad (45)$$

$$\log \pi(a|s) = \ell_a - \log \sum_{j=1}^3 \exp(\ell_j/\tau) \quad (46)$$

9.6 Mixture Output

The final prediction is a weighted combination of expert outputs:

$$p_{final} = \sum_{i=1}^3 \pi_i \cdot p_i \quad (47)$$

where π_i is the gating probability for expert i and p_i is the output probability from expert i .

9.7 Expert Routing

During training, each sample is routed to its selected expert:

$$p_{selected} = \begin{cases} p_1 & \text{if } a = 1 \\ p_2 & \text{if } a = 2 \\ p_3 & \text{if } a = 3 \end{cases} \quad (48)$$

During evaluation, the mixture output is used for consistent performance measurement.

10 Training Algorithm

10.1 Two-Phase Training

The MoE system uses a two-phase training approach:

10.1.1 Phase 1: Expert Pretraining

Each expert is pretrained independently for $\frac{T}{2}$ epochs:

$$\mathcal{L}_{total} = \mathcal{L}_{E1} + \mathcal{L}_{E2} + \mathcal{L}_{E3} \quad (49)$$

$$\mathcal{L}_{E1} = \text{CrossEntropy}(p_1, y) \quad (50)$$

$$\mathcal{L}_{E2} = \text{CrossEntropy}(p_2, y) + \lambda_{rep}\mathcal{L}_{rep} + \lambda_{fair}\mathcal{L}_{fair} \quad (51)$$

$$\mathcal{L}_{E3} = \text{CrossEntropy}(p_3, y) + \lambda_{attention}\mathcal{L}_{attention} + \lambda_{adv}\mathcal{L}_{adv} \quad (52)$$

10.1.2 Phase 2: Gate Training

The gating network is trained using REINFORCE for $\frac{T}{2}$ epochs.

10.2 REINFORCE Algorithm

The gate training uses policy gradient with the following components:

10.2.1 Reward Formulation

The reward function balances utility improvement and fairness improvement:

$$R = (U_2 - U_1) - (F_2 - F_1) \quad (53)$$

where:

$$U = \frac{\text{Accuracy} + \text{F1} + \text{AUC}}{3} \quad (54)$$

$$F = \frac{F_{\text{result}} + F_{\text{procedure}}}{2} \quad (55)$$

$$F_{\text{result}} = \frac{\text{DP} + \text{EO}}{2} \quad (56)$$

$$F_{\text{procedure}} = \frac{\text{REF} + \text{VEF} + \text{ATT}}{3} \quad (57)$$

10.2.2 Baseline and Advantage

A moving average baseline is used to reduce variance:

$$b_{t+1} = \alpha \cdot b_t + (1 - \alpha) \cdot R_t \quad (58)$$

$$A_t = R_t - b_t \quad (59)$$

10.2.3 Policy Loss

The policy loss includes entropy regularization and load balancing:

$$\mathcal{L}_{\text{gate}} = -\mathbb{E}[\log \pi(a|s) \cdot A] - \lambda_{\text{entropy}} \mathcal{H}(\pi) + \lambda_{\text{lb}} \mathcal{L}_{\text{lb}} \quad (60)$$

10.2.4 Entropy Regularization

The entropy term encourages exploration:

$$\mathcal{H}(\pi) = -\sum_{i=1}^3 \pi_i \log(\pi_i + \epsilon) \quad (61)$$

where ϵ is a small constant to avoid numerical issues.

10.3 Load Balancing

To ensure balanced expert usage, a KL divergence penalty is applied:

$$q_i = \frac{\text{usage_ma}_i + \epsilon}{\sum_{j=1}^3 (\text{usage_ma}_j + \epsilon)} \quad (62)$$

$$\mathcal{L}_{\text{lb}} = \text{KL}(q \parallel \text{Uniform}(1/3)) \quad (63)$$

where usage_ma_i is updated using exponential moving average:

$$\text{usage_ma}_i^{(t+1)} = \beta \cdot \text{usage_ma}_i^{(t)} + (1 - \beta) \cdot \pi_i^{(t)} \quad (64)$$

10.4 Expert Fine-tuning

During gate training, experts are periodically fine-tuned to maintain their performance:

$$\mathcal{L}_{\text{fine-tune}} = \mathcal{L}_{E1} + \mathcal{L}_{E2} + \mathcal{L}_{E3} \quad (65)$$

This fine-tuning occurs every 100 epochs to prevent expert degradation during gate training.

11 Evaluation Metrics

11.1 Utility Metrics

$$\text{Accuracy} = \frac{1}{n} \sum_{i=1}^n \mathbb{I}[\hat{y}_i = y_i] \quad (66)$$

$$\text{F1} = \frac{2 \cdot \text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}} \quad (67)$$

$$\text{AUC} = \int_0^1 \text{TPR}(t) \cdot \text{FPR}'(t) dt \quad (68)$$

11.2 Result Fairness Metrics

$$\text{DP Gap} = |\mathbb{E}[\hat{y}|s = 0] - \mathbb{E}[\hat{y}|s = 1]| \quad (69)$$

$$\text{EO Gap} = |\mathbb{E}[\hat{y}|s = 0, y = 1] - \mathbb{E}[\hat{y}|s = 1, y = 1]| \quad (70)$$

11.3 Procedural Fairness Metrics

$$\text{REF} = \text{Representation Fairness} \quad (71)$$

$$\text{VEF} = \text{Value Fairness} \quad (72)$$

$$\text{ATT} = \text{Attention JSD between groups} \quad (73)$$

11.4 Final Score

The overall performance is evaluated using:

$$\text{Final Score} = U - F_{\text{result}} - F_{\text{procedure}} \quad (74)$$

12 Implementation Details

12.1 Caching System

The system supports expert and gate caching:

- Expert weights are cached after pretraining
- Gate weights are cached after training
- Cache directory: `weights/moe_experts/`

12.2 Optimization

- Expert optimizer: Adam with learning rate lr and weight decay
- Gate optimizer: Adam with learning rate $gate_lr$ and no weight decay
- Momentum for baseline: $\alpha = 0.9$
- Momentum for usage tracking: $\beta = 0.9$

12.3 Training Hyperparameters

- Total epochs: T (default: 200)
- Expert pretraining: $T/2$ epochs
- Gate training: $T/2$ epochs
- Learning rates: $lr = 10^{-3}$, $gate_lr = 10^{-3}$
- Regularization: $\lambda_{entropy} = 10^{-3}$, $\lambda_{lb} = 10^{-3}$

13 Conclusion

The three expert models provide complementary approaches to fairness in machine learning:

- **Expert 1:** Pure utility optimization without fairness constraints
- **Expert 2:** Result-driven fairness through representation alignment and statistical parity
- **Expert 3:** Procedural fairness through attention alignment and adversarial debiasing

The gating network uses REINFORCE to learn optimal expert selection, balancing utility and fairness improvements. This multi-expert framework allows for flexible combination of different fairness approaches, enabling the system to adapt to various fairness requirements and trade-offs between accuracy and fairness.